

```
import numpy as np
```

```
class Softmax(object):
```

```
    def __init__(self, dims=[10, 3073]):
        self.init_weights(dims=dims)
```

```
    def init_weights(self, dims):
        """
        Initializes the weight matrix of the Softmax classifier.
        Note that it has shape (C, D) where C is the number of
        classes and D is the feature size.
        """
        self.W = np.random.normal(size=dims) * 0.0001
```

```
    def loss(self, X, y):
        """
        Calculates the softmax loss.

        Inputs have dimension D, there are C classes, and we operate on minibatches
        of N examples.

        Inputs:
        - X: A numpy array of shape (N, D) containing a minibatch of data.
        - y: A numpy array of shape (N,) containing training labels; y[i] = c means
            that X[i] has label c, where 0 <= c < C.

        Returns a tuple of:
        - loss as single float
        """
```

```
    # Initialize the loss to zero.
    loss = 0.0
```

```
    # ===== #
    # YOUR CODE HERE:
    #   Calculate the normalized softmax loss. Store it as the variable loss.
    #   (That is, calculate the sum of the losses of all the training
    #   set margins, and then normalize the loss by the number of
    #   training examples.)
    # ===== #
```

```
N, D = X.shape
scores = np.dot(X, self.W.T)
for i in range(N):
    score_ith = scores[i,:]
    score_ith -= np.max(score_ith)
    current_score = score_ith[y[i]]
    loss += current_score - np.log(np.sum(np.exp(score_ith)))
```

```
    # ===== #
    # END YOUR CODE HERE
    # ===== #
```

```
    return -loss/N
```

```
    def loss_and_grad(self, X, y):
        """
        Same as self.loss(X, y), except that it also returns the gradient.

        Output: grad -- a matrix of the same dimensions as W containing
            the gradient of the loss with respect to W.
        """
```

```
    # Initialize the loss and gradient to zero.
    loss = 0.0
```

```

grad = np.zeros_like(self.W)

# ===== #
# YOUR CODE HERE:
# Calculate the softmax loss and the gradient. Store the gradient
# as the variable grad.
# ===== #
N = X.shape[0]
C = self.W.shape[0]
scores = np.dot(X, self.W.T)
for i in range(N):
    score_ith = scores[i,:]
    score_ith -= np.max(score_ith)
    grad[y[i]] += X[i]
    for c in range(C):
        grad[c] -= np.exp(score_ith[c])/np.sum(np.exp(score_ith))*X[i]

loss = self.loss(X, y)
# ===== #
# END YOUR CODE HERE
# ===== #

return loss, -grad/N

def grad_check_sparse(self, X, y, your_grad, num_checks=10, h=1e-5):
    """
    sample a few random elements and only return numerical
    in these dimensions.
    """

    for i in np.arange(num_checks):
        ix = tuple([np.random.randint(m) for m in self.W.shape])

        oldval = self.W[ix]
        self.W[ix] = oldval + h # increment by h
        fxph = self.loss(X, y)
        self.W[ix] = oldval - h # decrement by h
        fxmh = self.loss(X,y) # evaluate f(x - h)
        self.W[ix] = oldval # reset

        grad_numerical = (fxph - fxmh) / (2 * h)
        grad_analytic = your_grad[ix]
        rel_error = abs(grad_numerical - grad_analytic) / (abs(grad_numerical) +
abs(grad_analytic))
        print('numerical: %f analytic: %f, relative error: %e' % (grad_numerical, grad_analytic,
rel_error))

def fast_loss_and_grad(self, X, y):
    """
    A vectorized implementation of loss_and_grad. It shares the same
    inputs and ouputs as loss_and_grad.
    """
    loss = 0.0
    grad = np.zeros(self.W.shape) # initialize the gradient as zero

    # ===== #
    # YOUR CODE HERE:
    # Calculate the softmax loss and gradient WITHOUT any for loops.
    # ===== #
    X = np.array(X)
    N = X.shape[0]
    C = self.W.shape[0]
    scores = np.dot(X, self.W.T)
    grad = np.sum(np.exp)
    scores = np.dot(X, self.W.T)
    scores = (scores.T - np.max(scores, axis=1)).T

```

```

    loss = -1/N * np.sum((scores[range(N), y] - np.log(np.sum(np.exp(scores), axis=1))),
axis=0)

```

```

probs = -np.exp(scores)/np.sum(np.exp(scores), axis=1, keepdims=True)
probs[range(N), y] += 1
grad = -1/N * np.dot(probs.T, X)

```

```

# ===== #
# END YOUR CODE HERE
# ===== #

```

```

return loss, grad

```

```

def train(self, X, y, learning_rate=1e-3, num_iters=100,
        batch_size=200, verbose=False):

```

```

    """

```

```

    Train this linear classifier using stochastic gradient descent.

```

```

    Inputs:

```

- *X*: A numpy array of shape (N, D) containing training data; there are N training samples each of dimension D.
- *y*: A numpy array of shape (N,) containing training labels; *y*[*i*] = *c* means that *X*[*i*] has label  $0 \leq c < C$  for *C* classes.
- *learning\_rate*: (float) learning rate for optimization.
- *num\_iters*: (integer) number of steps to take when optimizing
- *batch\_size*: (integer) number of training examples to use at each step.
- *verbose*: (boolean) If true, print progress during optimization.

```

    Outputs:

```

```

    A list containing the value of the loss function at each training iteration.

```

```

    """

```

```

    num_train, dim = X.shape
    num_classes = np.max(y) + 1 # assume y takes values 0...K-1 where K is number of classes

```

```

    self.init_weights(dims=[np.max(y) + 1, X.shape[1]]) # initializes the weights of self.W

```

```

    # Run stochastic gradient descent to optimize W

```

```

    loss_history = []

```

```

    for it in np.arange(num_iters):

```

```

        X_batch = None

```

```

        y_batch = None

```

```

        # ===== #
        # YOUR CODE HERE:

```

```

        # Sample batch_size elements from the training data for use in
        # gradient descent. After sampling,

```

```

        # - X_batch should have shape: (batch_size, dim)

```

```

        # - y_batch should have shape: (batch_size,)

```

```

        # The indices should be randomly generated to reduce correlations
        # in the dataset. Use np.random.choice. It's okay to sample with
        # replacement.

```

```

        # ===== #

```

```

        indexes = np.random.choice(num_train, batch_size)

```

```

        X_batch = X[indexes]

```

```

        y_batch = y[indexes]

```

```

        # ===== #

```

```

        # END YOUR CODE HERE

```

```

        # ===== #

```

```

        # evaluate loss and gradient

```

```

        loss, grad = self.fast_loss_and_grad(X_batch, y_batch)

```

```

        loss_history.append(loss)

```

```

        # ===== #

```

```

        # YOUR CODE HERE:

```

```

        # Update the parameters, self.W, with a gradient step

```

```

# ===== #
self.W -= learning_rate* grad

# ===== #
# END YOUR CODE HERE
# ===== #

if verbose and it % 100 == 0:
    print('iteration {} / {}: loss {}'.format(it, num_iters, loss))

return loss_history

def predict(self, X):
    """
    Inputs:
    - X: N x D array of training data. Each row is a D-dimensional point.

    Returns:
    - y_pred: Predicted labels for the data in X. y_pred is a 1-dimensional
      array of length N, and each element is an integer giving the predicted
      class.
    """
    y_pred = np.zeros(X.shape[1])
    # ===== #
    # YOUR CODE HERE:
    #   Predict the labels given the training data.
    # ===== #
    # predicted = X.dot(self.W.T)
    # #normalize
    # predicted = predicted - np.argmax(predicted)

    # y_predict = np.exp(predicted) / np.sum(np.exp(predicted))
    # y_pred = np.argmax(y_predict)

    scores = np.dot(X, self.W.T) # scores.shape = num_examples * num_classes
    scores = (scores.T - np.max(scores, axis=1)).T # to control the overflow

    probs = np.exp(scores) / np.sum(np.exp(scores), axis=1, keepdims=True)
    y_pred = np.argmax(probs, axis=1)

    # ===== #
    # END YOUR CODE HERE
    # ===== #

    return y_pred

```